

Due Friday 16th of June at 6pm Sydney time (week 3)

In this assignment we review some basic algorithms and data structures, and we apply the divide-and-conquer paradigm. There are *three problems* each worth 20 marks, for a total of 60 marks. Partial credit will be awarded for progress towards a solution. We'll award one mark for a response of "one sympathy mark please" for a whole question, but not for parts of a question.

Any requests for clarification of the assignment questions should be submitted using the [Ed forum](#). We will maintain a [FAQ thread](#) for this assignment.

For each question requiring you to design an algorithm, you *must* justify the correctness of your algorithm. If a time bound is specified in the question, you also *must* argue that your algorithm meets this time bound. The required time bound always applies to the *worst case* unless otherwise specified.

You must submit your response to each question as a separate PDF document on Moodle. You can submit as many times as you like. Only the last submission will be marked.

Your solutions must be typed, *not* handwritten. We recommend that you use LaTeX, since:

- as a UNSW student, you have a free Professional account on [Overleaf](#), and
- we will release a LaTeX template for each assignment question.

Other typesetting systems that support mathematical notation (such as Microsoft Word) are also acceptable.

Your assignment submissions must be your own work.

- You may make reference to published course material (e.g. lecture slides, tutorial solutions) without providing a formal citation. The same applies to material from COMP2521/9024.
- You may make reference to either of the recommended textbooks with a citation in any format.
- You may reproduce general material from external sources in your own words, along with a citation in any format. 'General' here excludes material directly concerning the assignment question. For example, you can use material which gives more detail on certain properties of a data structure, but you cannot use material which directly answers the particular question asked in the assignment.
- You may discuss the assignment problems privately with other students. If you do so, you must acknowledge the other students by name and zID in a citation.
- However, you must write your submissions entirely by yourself.
 - Do not share your written work with anyone except COMP3121/9101 staff, and do not store it in a publicly accessible repository.
 - The only exception here is [UNSW Smarthinking](#), which is the university's official writing support service.

Please review the UNSW policy on [plagiarism](#). Academic misconduct carries severe penalties.

Please read the [Frequently Asked Questions](#) document, which contains extensive information about these assignments, including:

- how to get help with assignment problems, and what level of help course staff can give you
- extensions, Special Consideration and late submissions
- an overview of our marking procedures and marking guidelines
- how to appeal your mark, should you wish to do so.

Question 1

You are given an array A of n distinct positive integers and a positive integer x .

1.1 [8 marks] Design an algorithm which decides if there exist two distinct indices $1 \leq i, j \leq n$ such that $2A[i] - 3A[j] = x$. In the worst-case, your algorithm must run in $O(n \log n)$ time.

1.2 [4 marks] Solve the same problem as in 1.1 but with an algorithm which runs in the **expected time** of $O(n)$.

1.3 [8 marks] Design an algorithm which counts how many distinct pairs of indices (i, j) where $1 \leq i < j \leq n$ satisfy both:

- $A[i] > A[j]$; and
- $A[i] + A[j] = x$.

In the worst-case, your algorithm must run in $O(n(\log n)^2)$ time.

Question 2

2.1 [6 marks] Blake and Red each have a collection of n distinct Pokemon cards. Blake creates an array B containing the serial numbers of their cards, and Red creates an array R with the serial numbers of his. Blake wishes to know how many of their cards are *not* also owned by Red (that is, the size of $B \setminus R$).

For example, if Blake has cards with serial numbers $B = [7, 2, 1, 5]$ and Red has cards numbered $R = [2, 8, 7, 3]$, then there are 2 cards owned by Blake but not Red.

Design an algorithm which runs in $O(n \log n)$ time and determines how many cards are owned by Blake but not Red.

2.2 [4 marks] Gerald has a class of k students, each of which has a collection of distinct Pokemon cards. One day, he asks each member of his class to bring their collection, sorted by serial number. The i th student has c_i cards, with their **sorted** serial numbers in an array $N_i[1 \dots c_i]$. The total number of (not necessarily distinct) cards owned by the class is $S = \sum_{i=1}^k c_i$.

He wants to find all the distinct cards owned by the class (i.e. $\bigcup_{i=1}^k N_i$), but needs to do so in $O(S \log k)$ to finish the lesson in time. Gerald devises the following algorithm to do so:

We create an array N^* to store the serial numbers of the unique cards owned by the class. We first put all the numbers from N_1 into N^* .

Then, for each i from 2 to k , we merge N_i with N^* , only including one copy of duplicate cards. Since N_i and N^* are both sorted, we are guaranteed to see duplicate cards one after the other while merging. We then replace N^* with the merged array.

The final array N^* contains the merged arrays N_1, N_2, \dots, N_k without duplicate serial numbers, i.e. all the distinct cards owned by the class.

This algorithm is correct, but unfortunately does not meet Gerald's $O(S \log k)$ time complexity requirement. Justify why the worst-case time complexity of Gerald's algorithm is slower than $O(S \log k)$.

Note that an example with a fixed size (i.e., fixed S or k) is **not** sufficient!

2.3 [10 marks] Design an algorithm that finds the distinct cards owned by Gerald's class, and runs in $O(S \log k)$ time.

Question 3

3.1 Read about the asymptotic notation in the review material and determine if $f(n) = O(g(n))$ or $g(n) = O(f(n))$ or both (i.e., $f(n) = \Theta(g(n))$) or neither of the two, for the following pairs of functions.

[A] [5 marks]

$$f(n) = \log_2(n); \quad g(n) = \sqrt[5]{n}$$

[B] [5 marks]

$$f(n) = n^n; \quad g(n) = 2^{n \log_2(n^2)}$$

[C] [5 marks]

$$f(n) = n^{1+\cos(\pi n)}; \quad g(n) = n$$

You might find L'Hôpital's rule useful: if $f(x), g(x) \rightarrow \infty$ as $x \rightarrow \infty$ and they are differentiable, then

$$\lim_{x \rightarrow \infty} \frac{f(x)}{g(x)} = \lim_{x \rightarrow \infty} \frac{f'(x)}{g'(x)}.$$

3.2 [5 marks] Given n strings, each of length at most m , a divide and conquer approach can be used to find the longest common prefix amongst the strings. For example, the longest common prefix amongst **apple**, **apply** and **apart** is **ap**. Song has provided an algorithm to do so below:

1. Recursively determine the longest common prefixes among the first $n/2$ words and the last $n/2$ words.
2. Merge the results to determine the longest common prefix between the two halves by scanning through character by character.
3. Base case: for $n = 1$, the longest common prefix is the entire string.

However, Song isn't sure about the time complexity of the above algorithm. In Big-Theta notation, explain and justify what the time complexity of the above algorithm is.